

Software Security 05

Sven Schäge

Methods for Finding Security Vulnerabilities

Sven Schäge

The Complementary 3

- There are three main methods to find security vulnerabilities in practice
 - Static Code Analysis
 - At compile time
 - Fuzzing (smart generation of tests and subsequent testing of software)
 - At runtime
 - Review by experts (penetration testing)
 - Formal Verification?
- Each method has benefits and drawbacks; however, they complement each other nicely
- Nevertheless, they cannot guarantee the absence of security vulnerabilities

Static Code Analysis, Static Application Security Testing (SAST)

- + Automatic
 - Most of the work is done by tools that automatically scan the code
 - Scales well
- + Computationally cheap
- + Does not need the software to be run
- - Can raise false alarms
 - Not guaranteed to find all bugs – this depends on the specification of bug
- - Cannot find bugs that only reveal themselves when software is running in complex runtime environments
- + Cover the entire code even parts that will be executed only in rare cases
- - Cannot be used to find errors in runtime environment
- + Finding responsible code positions is easy

Fuzzing and Dynamic Application Security Testing (DAST)

- + Automatic
 - Scales well
- + Cannot raise false alarms
- - Computationally more expensive
 - May scan year-round
- - Needs the software to be run (dynamic analysis)
- - Not guaranteed to find all bugs
- + Can find security vulnerabilities in deployed environments
- + Does not require the code
- - Testing may launch an attack!
 - Testing should be performed in production-like but non-production environment.
- - Cannot cover all of the source code and thus application
- + Can find errors in runtime environment
- + Can be used to verify results of Static Code Analysis
- - After finding a vulnerability, we have to find the code position that is responsible for it

Review by Experts

- - Manual inspection with focus on security
 - Does not scale well
- - Requires time and can be (comparably) expensive
- + No false alarms
- - Not guaranteed to find all bugs

Important Quality Metrics

Completeness and Soundness of Analysis

- When classifying programs, it depends on the mandate (the property that programs should be classified by) used whether the classifier can be called complete or sound.
- In this lecture we will restrict ourselves to the following understanding of completeness and soundness **unless the mandate is not explicitly stated**.
 - If we say a classifier/program checker/analysis tool is complete we mean that it will classify all programs that have bugs as having bugs. (Minimizes false negatives, ensuring that the tool does not miss real vulnerabilities)
 - If we say that a classifier/program checker/analysis tool is sound we mean that it will always classify programs that are bug-free as bug-free. Soundness in software security refers to the assurance that a security tool correctly identifies all actual security vulnerabilities without reporting false positives.
 - This means we use the mandate **bugginess**, trying to answer the question „Is the program buggy?“
- In general, inverting the mandate of a classifier will swap the results with respect to soundness and completeness!

Measuring Accuracy: Precision and Recall for Finding Buggy Programs

the terms in the table are buildt as X Y where
X ∈ {True, False} indicates correct or incorrect classification and
Y ∈ {Positive, Negative} indicates the result of the Analysis

Program Ground Truth		Analysis Outcome (Marked as having no security weakness)	
		Alarm	No Alarm
Buggy	Bug-free	True Positive	False Negatives
		False Positive	True Negative

program is rightfully declared to contain a bug

buggy programs that have not been caught

false alarms

program is rightfully declared to not contain a bug

Completeness and Soundness of Analysis

- For our purposes, (perfect) completeness means if a program is claimed to be buggy then it always is buggy.
- (Perfect) Soundness means that if a programs is claimed to be buggy the it is with buggy will always be identified as buggy.
- Precision and Recall quantify these absolute categories
- Precision
 - Of the items flagged as buggy, how many are truly buggy?“
 - $\text{Precision} = \frac{\text{True Positives}}{(\text{False Positives} + \text{True Positives})}$
 - No False Positives (no false alarm) \Rightarrow Precision=1
 - 100 security alerts but only 90 of them are true positives it means that we have a precision of 90%
- Recall:
 - Of all the truly buggy items, how many did the system correctly identify?“
 - $\text{Recall} = \frac{\text{True Positives}}{(\text{True Positives} + \text{False Negatives})}$
 - No False Negatives (all bugs caught) \Rightarrow Recall=1
- **False Positives (FP):** Items that are incorrectly classified as buggy when they are not. (false alarm)
- **False Negatives (FN):** Items that are buggy but are not identified as such by the security tool. (bug that is not caught)
- **True Positives (TP):** Items that are correctly identified as buggy.
- **True Negatives (TN):** Items that are correctly identified as bug-free.

F-Measure

- The F-measure is a standard measure of accuracy, combining precision and recall
- Harmonic mean of recall and precision
- Hypothetical, ideal analysis: F-measure=1
- Here precision and recall are equally important

$$\begin{aligned}\text{F-measure} &= 2 / ((1/\text{precision}) + (1/\text{recall})) \\ &= 2 * \text{recall} * \text{precision} / (\text{recall} + \text{precision})\end{aligned}$$

Fx-Measure

- Depending on the applications we may want to weigh precision and recall differently
- The Fx-measure accounts for this using parameter x
- The Fx-measure was derived so that it measures the effectiveness with respect to a user who attaches x times as much importance to recall as precision
- Two commonly used values for x are 2, which weighs recall higher than precision, and 0.5, which weighs recall lower than precision.

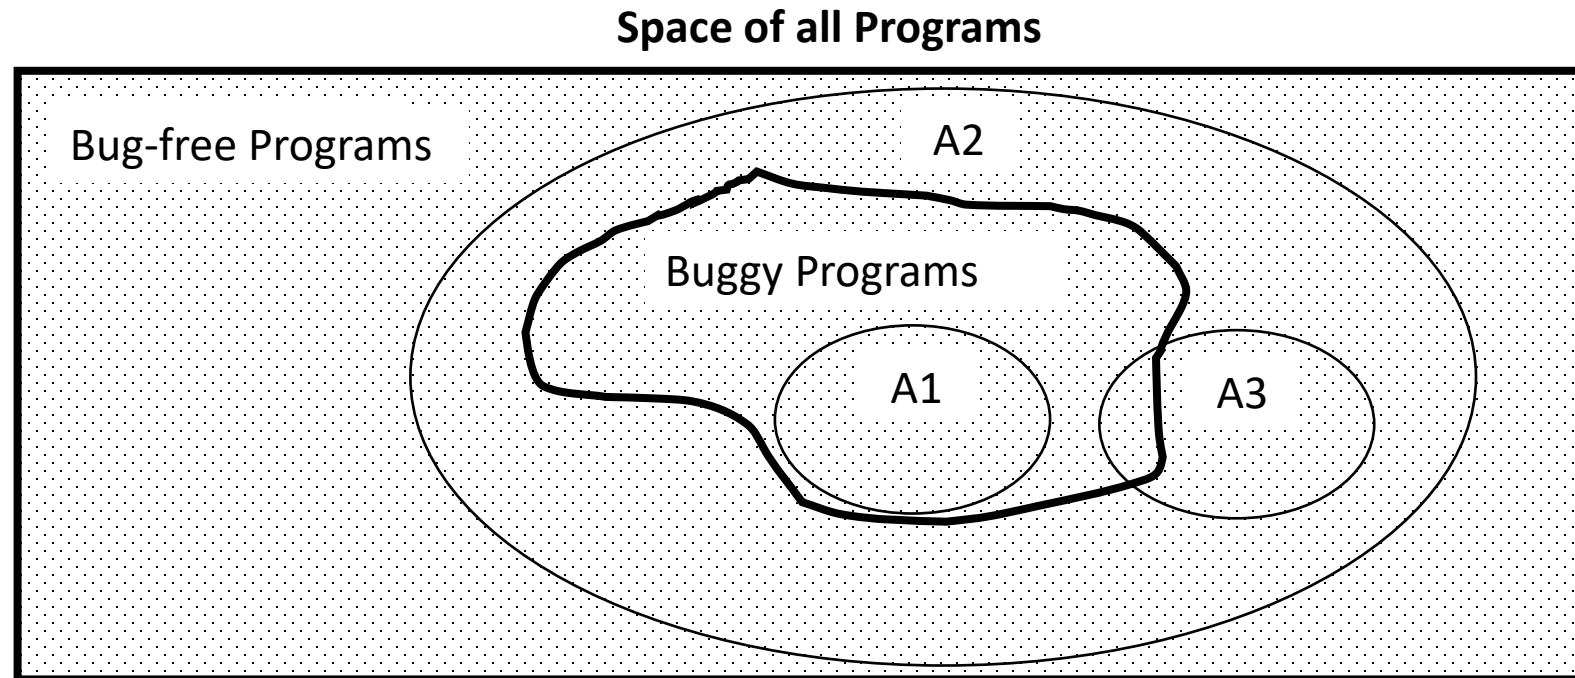
$$\text{Fx-measure} = (1+x^2) * \text{recall} * \text{precision} / (\text{recall} + (x^2 * \text{precision}))$$

Application-Aware Choice of Analysis

Choosing an Appropriate Program Analysis: Completeness vs. Soundness

- Assume the analysis tools A1, A2, A3 recognize their respective inner spaces as buggy programs, their outer space as bug-free programs. Assume their mandate is to find buggy programs.

	Complete?	Sound?
A1	Yes	No
A2	No	Yes
A3	No	No



On Specifications and Properties

Analysis and Specification

- Usually software analysis depends on a specification of some desirable property of a program
 - These specifications can vary in formality
 - More formal specifications are better for automatic testing
- For example, if we want to avoid division by zero, we could define a simple state machine to capture this as a specification.

Classifications of Specifications

- History-based specification
 - behavior based on system histories
 - assertions are interpreted over time
- State-based specification
 - behavior based on system states
 - series of sequential steps, (e.g. a financial transaction)
- Transition-based specification
 - behavior based on transitions from state-to-state of the system
- Functional specification
 - specify a system as a structure of mathematical functions
- Operational Specification
 - e.g. Petrinets, Algebras

Classification of Properties

- Safety Properties
 - Program will never reach a bad state
 - Examples:
 - assertions
 - predicate at point in program that should always evaluate to true
 - types and type system
 - define how variables should be interpreted, e.g. int
 - pre- and post conditions
 - a precondition is a condition or predicate that must always be true just prior to the execution of some section of code
 - a postcondition is a condition or predicate that must always be true just after the execution of some section of code
 - loop and class invariants
- Liveness Properties
 - Program will eventually reach a bad state
 - Examples:
 - program termination,
 - starvation freedom

Methods for Finding Security Vulnerabilities – Static Code Analysis

Sven Schäge

What can we Hope to Achieve?

- Halting Problem and Rice's Theorem
There is no algorithm that inputs any program and decides if the program satisfies a given non-trivial semantic property.
- Wishlist
 - Complete (no uncaught bugs)
 - Sound (no false alarms)
 - Automatic
 - Powerful/Non-trivial
- => We must give up at least one point on the wishlist!

Static Code Analysis, Static Application Security Testing (SAST)

- + Automatic
 - Most of the work is done by tools that automatically scan the code
 - Scales well
- + Computationally cheap
- + Does not need the software to be run
- - Can raise false alarms
- - Not guaranteed to find all bugs
- - Cannot find bugs that only reveal themselves when software is running in complex runtime environments
- + Cover the entire code even parts that will be executed in extreme cases
- - Cannot be used to find errors in runtime environment
- + Finding responsible code positions is easy

Static Code Analysis

- Mainly, analyse source code
- Often concentrates on syntactic information
 - Easier to analyse automatically
- Search for patterns or code fragments that meet some rule
 - e.g. every pointer that has been given memory needs to be freed
- Several degrees of locality
 - Consider only single statement at hand
 - or context in entire code
- In general, help the developer to get additional information and new perspectives on the code often in the form of:
 - at position X in the control flow graph we have the property Y

Concepts and Techniques for Static Code Analysis

- Formalization and Visualizations (of code dependencies)
 - Control Flow Analysis (Call Graphs and Control Flow Graphs)
- Iterative Information Generation
 - Data Flow Analysis
- Pattern checking in code for compliance with coding rules
- Type Systems
 - Most widespread form of static analysis
 - Helps to avoid that variables are combined in a way that makes no sense: e.g. add int and char
- Instrumentation of binary code
 - allow for better analysis of program behaviour
 - and trace back origin of problems in source code
 - May add additional checks to the binary that can be evaluated when running the software
- Strong overlap of techniques with techniques used to build compilers

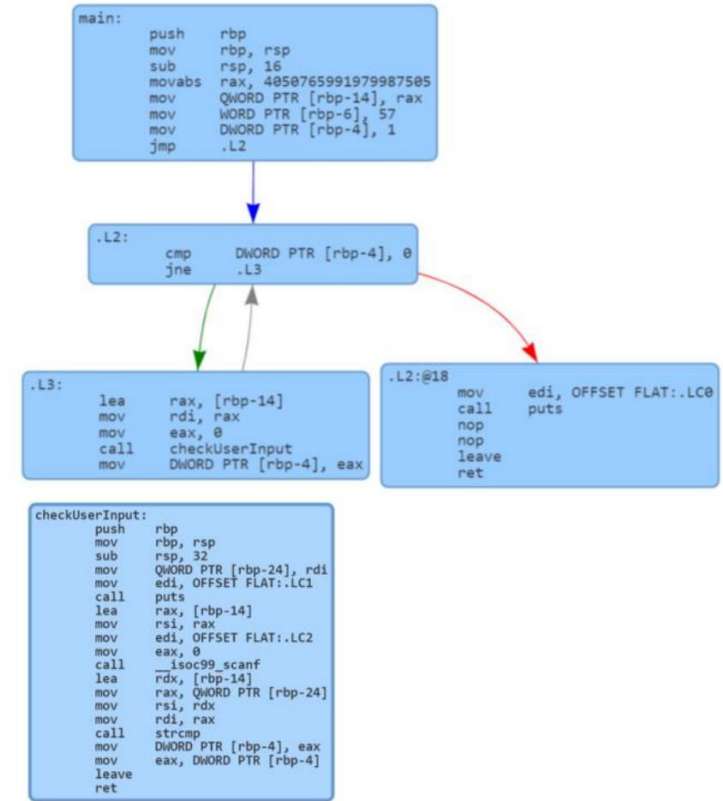
Call Graphs (CG)

- Graph representation of all function calls that a program could make
 - Each node represents a function
 - Edges represent function calls
- Calling relationships between subroutines in a program
- Inter-procedural view of program

Control Flow Graphs (CFG)

- Graph representation of all paths that can/might be traversed through a function during its execution
 - Each node represents a basic block, i.e. straight-line code fragment without any jumps or jump targets
 - Jump targets start a block, jumps end a block
- Generation of Control Flow Graph
 - Starting from Full Flow Graph (every instruction is represented by a node) via edge contraction
- Intra-procedural view of program

Graph output



Prerequisite: The WHILE Language

```
x = 5;  
y = 1;  
while (x != 1) {  
    y = x * y;  
    x = x - 1  
}
```

(statement) $S ::= x = a \mid S1 ; S2 \mid$
 $\text{if } (b) \{ S1 \} \text{ else } \{ S2 \} \mid$
 $\text{while } (b) \{ S1 \}$



(arithmetic expression) $a ::= x \mid n \mid a1 * a2 \mid a1 - a2$

(boolean expression) $b ::= \text{true} \mid !b \mid b1 \ \&\& \ b2 \mid$
 $a1 \ != \ a2$


(integer variable) x

(integer constant) n

Control-Flow Graph



Control-Flow Graph



```
x = 5;
while (x != 0) {
    y = x;
    x = x - 1;
    while (y != 0) {
        y = y - 1
    }
}
```

Inter-Procedural Control Flow Graphs

- Inter-procedural Control Flow Graph
- While CFGs represent the control flow of a single procedure, Inter-Procedural Control Flow Graphs (ICFG) represent the control flow of entire programs
- Combination of CG and CFG

Static vs. Dynamic Call Graphs

- Static Call Graph (undecidable problem -> approximation only)
 - A static call graph is a call graph intended to represent every possible run of the program.
 - The exact static call graph is an undecidable problem, so static call graph algorithms are generally overapproximations.
 - That is, every call relationship that occurs is represented in the graph, and possibly also some call relationships that would never occur in actual runs of the program.
- Dynamic Call Graph (what parts of the code have been executed in a specific run)

Pattern Checking in Code for Compliance with Coding Rules/Policies

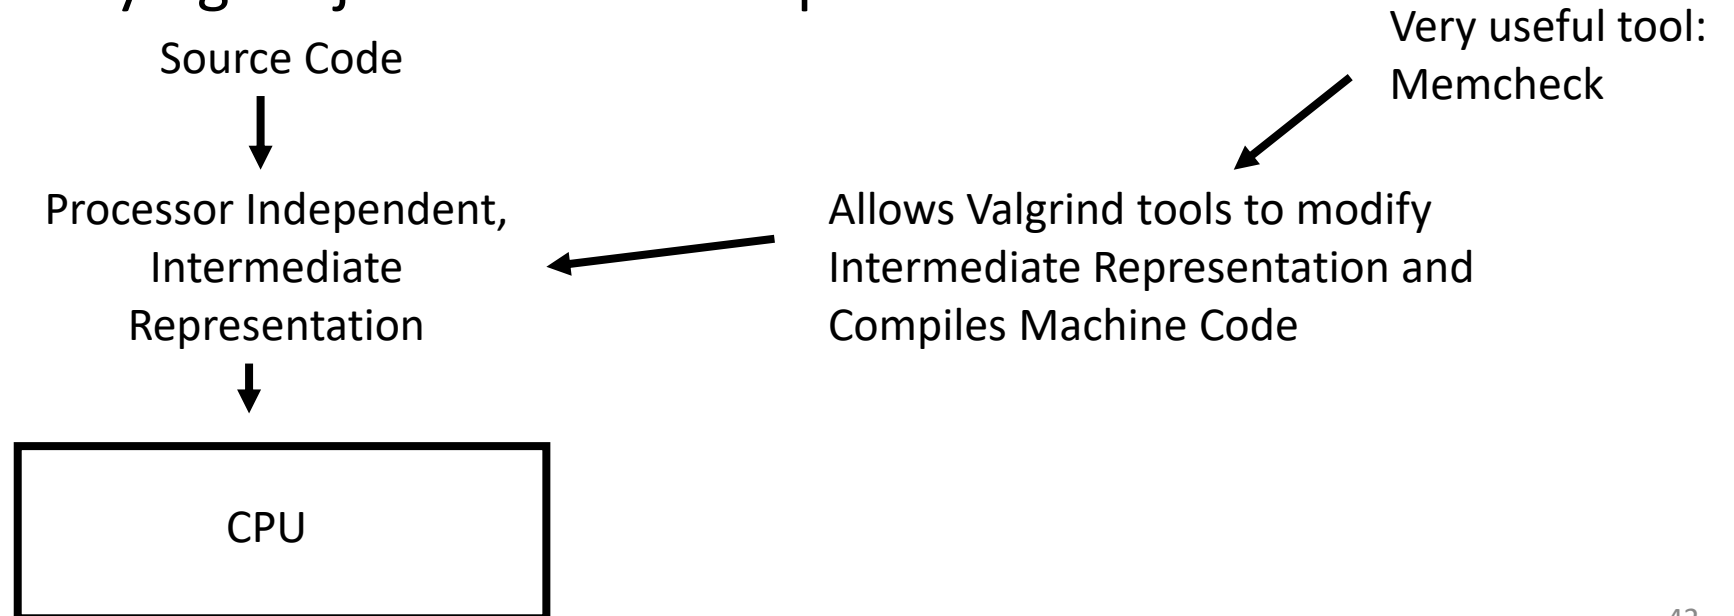
- Varying degrees of „quality“ of pattern checking (with varying Speed)
 - Only syntactical
 - Do not use function X from library Y
 - More context-sensitive
 - Number of allocations and free operations should balance
- Policies can be defined by company
- Policies can restrict to use of subset of language only
 - MISRA-C

Binary Code Instrumentation

- Adds additional information to binary code
 - Resulting binary much larger than stripped-off binary
 - Typically used for debugging
 - Allows to connect execution of binary to source code
 - Not shipped in practice to users as it may greatly help attacker, e.g. when trying to decompile (generate code from binary in highlevel language with equivalent functionality)
- General techniques for debugging (non-specific to security analysis)
 - Step-wise execution
 - Breakpoints, i.e. program execution halts at specific points
 - Watchpoints , i.e. if certain conditions are met, the debugger gives feedback and may halt the execution

Binary Code Instrumentation

- Add extra code to resulting binary
- May be used to search for errors that are often security critical
- Example: Valgrind
 - Tool for memory debugging, memory leak detection and profiling
- Virtual machine relying on just-in-time compilation



Valgrinds Memcheck

- Validity
 - All freshly allocated memory is not used until it is initialized first
- Addressability
 - Memory used in new allocations is indeed free
 - In contrast to non-freed memory
- Technically, replaces standard C memory allocator
 - Can keep track of state of memory and implement guards around memory areas to detect buffer overflows
- Memory leaks (allocated memory that is not freed), use-after free, buffer overflows, use of non-initialized memory
- Does not find all security vulnerabilities!
- Costs of Valgrind: program execution about 20-30 times slower
 - Usually, only used in testing phase

Other Concepts

- Check code for compliance with a specific subset of programming language that is expected to deliver better (security-wise) code
 - e.g. MISRA-C
- Advanced concepts: Runtime Monitoring
- Profiling run program and output additional information afterwards